

# Write up of ASG2: Multithreading and server caching

Yujia Li

CruzID: yli302

## Write up questions:

• Using either your HTTP client or curl and your original HTTP server from Assignment 1, do the following:

• Place four different large files on the server. This can be done simply by copying the files to the server's directory, and then running the server. The files should be around 4 MiB long.

• Start httpserver.

• Start four separate instances of the client at the same time, one GETting each of the files and measure (using time(1)) how long it takes to get the files. Perhaps the best way to do this is to write a simple shell script (command file) that starts four copies of the client program in the background, by using & at the end.

Filesize: 2288 KB, 3012 KB, 3442 KB, 8329 KB

1234567890123456789012345678901234567890	2019/1/25 11:27	文件	2,288 KB
1234567890123456789012345678901234567891	2018/12/13 15:17	文件	3,012 KB
1234567890123456789012345678901234567892	2018/8/26 19:54	文件	3,442 KB
1234567890123456789012345678901234567893	2018/10/20 14:27	文件	8,329 KB

Output:

```
yujia@yujia-VirtualBox:~/cmpe105/Assignment/test$ time (./httpclient 0.0.0.0:8000 r:123456789012345678901234567890:file0 & ./httpclient 0.0.0.0:8000 r:1234567890123456789012345678901234567891:file1 & ./httpclient 0.0.0.0:8000 r:1234567890123456789012345678901234567892:file2 & ./httpclient 0.0.0.0:8000 r:1234567890123456789012345678901234567893:file3)
-----
-----
-----
-----
-----
HTTP/1.1 200 OK
HTTP/1.1 200 OK
HTTP/1.1 200 OK
HTTP/1.1 200 OK

real    0m1.535s
user    0m0.031s
sys     0m0.221s
```

real 0m1.535s

user 0m0.031s

sys 0m0.221s

- Repeat the same experiment after you implement multi-threading. Is there any difference in performance?

Output:

```
yuji@yuji-VirtualBox:~/cmpe105/Assignment/test$ time (./httpclient 0.0.0.0:8000 r:123456789012345678901234567890:file0 & ./httpclient 0.0.0.0:8000 r:123456789012345678901234567890:file1 & ./httpclient 0.0.0.0:8000 r:123456789012345678901234567890:file2 & ./httpclient 0.0.0.0:8000 r:123456789012345678901234567890:file3)
-----
-----
HTTP/1.1 200 OK
-----
-----
HTTP/1.1 200 OK
-----
-----
HTTP/1.1 200 OK
-----
-----
real    0m0.528s
user    0m0.039s
sys     0m0.200s
```

real 0m0.528s

user 0m0.039s

sys 0m0.200s

Multi-threading is much faster about 1s than original HTTP server.

- What is likely to be the bottleneck in your system? How much concurrency is available in various parts, such as dispatch, worker, cache? Can you increase concurrency in any of these areas and, if so, how?

The bottleneck must be the synchronization about shared variables, especially about read/write operation. In my implementation the semaphore about cache will affect dispatch thread, so the performance will decrease. I lock read and write starting at open() in the work threads, so most of the read/write operation wasted many time and this will lead to dispatch thread waiting for the read and write operation finish. Performance can be increased by separating the concurrency of cache and work threads. For example, when cache write/read to/from disk with different httpname file, the work threads do not need to wait. But it is difficult for me, because I don't know how to figure whether working work threads is working on a same httpname file.

## TEST:

Multi-thread without cache can work with little bugs (see README)

Cache is not finished.



